

## Programming Exercise 4: Logistic Regression (with tensorflow)

=====

*# Run this cell to import all the libraries you need for this exercise*

```
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import random
```

*# STEP 1: Read in data*

```
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets('MNIST_data', one_hot=True)
```

```

def TRAIN_SIZE(num):
    print ('Total Training Images in Dataset = ' + str(mnist.train.images.shape))
    print ('-----')
    x_train = mnist.train.images[:num,:]
    print ('x_train Examples Loaded = ' + str(x_train.shape))
    y_train = mnist.train.labels[:num,:]
    print ('y_train Examples Loaded = ' + str(y_train.shape))
    print('')
    return x_train, y_train

def TEST_SIZE(num):
    print ('Total Test Examples in Dataset = ' + str(mnist.test.images.shape))
    print ('-----')
    x_test = mnist.test.images[:num,:]
    print ('x_test Examples Loaded = ' + str(x_test.shape))
    y_test = mnist.test.labels[:num,:]
    print ('y_test Examples Loaded = ' + str(y_test.shape))
    return x_test, y_test

def display_digit(num,x_train,y_train):
    print(y_train[num])
    label = y_train[num].argmax(axis=0)
    image = x_train[num].reshape([28,28])
    plt.title('Example: %d Label: %d' % (num, label))
    plt.imshow(image, cmap=plt.get_cmap('gray_r'))
    plt.show()

```

```


Total Training Images in Dataset = (55000, 784)
-----
x_train Examples Loaded = (5500, 784)
y_train Examples Loaded = (5500, 10)

Total Test Examples in Dataset = (10000, 784)
-----
x_test Examples Loaded = (1000, 784)
y_test Examples Loaded = (1000, 10)


```

```
# Display 30 random digits
for i in range(3):
    selected = random.randint(0, 5500-1)
    display_digit(selected, X_train, Y_train)
```

[0. 0. 0. 0. 0. 0. 0. 1. 0. 0.]



[0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]



```
# n features, k classes
```

```
n = 784
```

```
k = 10
```

```
# INSTRUCTIONS: Complete the code for Steps 3-7
```

```
# STEP 3: Create placeholders for features and labels
```

```
# ===== YOUR CODE =====
```

```
#
```

```
# STEP 4: Create weights and bias
```

```
# ===== YOUR CODE =====
```

```
#
```

```
# STEP 5: Predict Y from X and W, b
```

```
# ===== YOUR CODE =====
```

```
#
```

```
# STEP 6: Define loss function
```

```
# ===== YOUR CODE =====
```

```
#
```

```
# STEP 7: Gradient Descent
```

```
# ===== YOUR CODE =====
```

```
#
```

```
# STEP 8: Calculate accuracy
```

$$P(Y = i|x, W, b) = \text{softmax}_i(Wx + b) \\ = \frac{e^{W_i x + b_i}}{\sum_j e^{W_j x + b_j}}$$

$$y_{pred} = \text{argmax}_i P(Y = i|x, W, b)$$

GradientDescentOptimizer

```
# Initializing the variables
```

```
init = tf.global_variables_initializer()
```

```
# Launch the graph
```

```
with tf.Session() as sess:  
    sess.run(init)
```

```
# Train model
```

```
for i in range(train_steps+1):
```

```
    sess.run([optimizer, loss], feed_dict={X: X_train, Y: Y_train})
```

```
    if i%100 == 0:
```

```
        print('Training Step: ', i)
```

```
        print('Accuracy:', sess.run(accuracy, feed_dict={X: X_test
```

```
# Obtain weights
```

```
W_values = sess.run(w)
```

```
Training Step: 1300
```

```
Accuracy: [ 0.851 85.1 ]
```

```
Training Step: 1400
```

```
Accuracy: [ 0.855 85.5 ]
```

```
Training Step: 1500
```

```
Accuracy: [ 0.855 85.5 ]
```

```
Training Step: 1600
```

```
Accuracy: [ 0.855 85.5 ]
```

```
Training Step: 1700
```

```
Accuracy: [ 0.856 85.6 ]
```

```
Training Step: 1800
```

```
Accuracy: [ 0.857 85.7 ]
```

```
Training Step: 1900
```

```
Accuracy: [ 0.858 85.8 ]
```

```
Training Step: 2000
```

```
Accuracy: [ 0.859 85.9 ]
```

```
Training Step: 2100
```

```
Accuracy: [ 0.86 86. ]
```

```
Training Step: 2200
```

```
Accuracy: [ 0.86 86. ]
```

```
Training Step: 2300
```

```
Accuracy: [ 0.863 86.3 ]
```


```
Training Step: 2400
```

```
Accuracy: [ 0.864 86.4 ]
```


```
Training Step: 2500
```

```
Accuracy: [ 0.865 86.5 ]
```


0




1




2



3



4



# EXERCISE 5



## Programming Exercise 5: Neural Networks (without tensorflow)

=====

In this exercise, we are going to build a neural network with

- 1) an input layer matching the size of our instance data (400 + the bias unit),
- 2) a hidden layer with 25 units (26 with the bias unit), and
- 3) an output layer with 10 units corresponding to our one-hot encoding for the class labels.

*# Run this cell to import all the libraries you need for this exercise*

```
import numpy as np
import matplotlib.pyplot as plt
from scipy.io import loadmat
```

*# STEP 1: Read in data*

```
data = loadmat('ex5data1.mat')
```


*# Assuming m examples...*

*# X should be a (m x 400)-matrix, each row containing the 400 features of an example*

*# and y should be a (m x 1)-matrix containing the labels associated with the m examples*

```
X = data['X']
y = data['y']
print(X.shape, y.shape)
```

```
(5000, 400) (5000, 1)
```





```

# standard sigmoid function
def sigmoid(z):
    return 1 / (1 + np.exp(-z))

# Forward propagation
# Assuming:
#   input_size = 400 (number of neurons in input layer, also number of features for each example)
#   hidden_size = 25 (number of neurons in hidden layer)
#   num_labels = 10 (number of neurons in output layer)
# Inputs of the function:
#   X          (m x 400)-matrix   containing the input examples
#   theta1     (25 x 401)-matrix  containing the weights between the input layer of 401 neurons and the hidden layer
#   theta2     (10 x 26)-matrix   containing the weights between the hidden layer of 25 neurons and the output layer
# Outputs of the function
#   a1         (m x 401)-matrix   essentially X with an additional column of 1s inserted as column 0
#   z2         (m x 25)-matrix   contains a1 and theta1 multiplied together
#   a2         (m x 26)-matrix   apply sigmoid function to z2 with an additional column of 1s inserted as column 0
#   z3         (m x 10)-matrix   contains a2 and theta2 multiplied together
#   h          (m x 10)-matrix   contains the result of applying sigmoid to z3
def forward_propagate(X, theta1, theta2):
    m = X.shape[0]
    a1 =    ### FILL-IN |
    z2 =    ### FILL-IN
    a2 =    ### FILL-IN
    z3 =    ### FILL-IN
    h =     ### FILL-IN
    return a1, z2, a2, z3, h

```

```

# Cost function
def cost(params, input_size, hidden_size, num_labels, X, y, learning_rate):
    m = X.shape[0]
    X = np.matrix(X)
    y = np.matrix(y)

    # reshape the parameter array into parameter matrices for each layer
    theta1 = np.matrix(np.reshape(params[:hidden_size * (input_size + 1)], (hidden_size, (input_size + 1))))
    theta2 = np.matrix(np.reshape(params[hidden_size * (input_size + 1):], (num_labels, (hidden_size + 1))))

    # run the feed-forward pass
    a1, z2, a2, z3, h = forward_propagate(X, theta1, theta2)

    # compute the cost
    J = 0
    for i in range(m):
        first_term = np.multiply(-y[i,:], np.log(h[i,:]))
        second_term = np.multiply((1 - y[i,:]), np.log(1 - h[i,:]))
        J += np.sum(first_term - second_term)

    J = J / m

    # Add regularization term
    J += (float(learning_rate) / (2 * m)) * (np.sum(np.power(theta1[:,1:], 2)) + np.sum(np.power(theta2[:,1:], 2)))

    return J

```

```

# initial setup
input_size = 400
hidden_size = 25
num_labels = 10
learning_rate = 1

# randomly initialize a parameter array of the size of the full network's parameters
params = (np.random.random(size=hidden_size * (input_size + 1) + num_labels * (hidden_size + 1)) - 0.5) * 0.25

m = X.shape[0]
X = np.matrix(X)
y = np.matrix(y)

# unravel the parameter array into parameter matrices for each layer
theta1 = np.matrix(np.reshape(params[:hidden_size * (input_size + 1)], (hidden_size, (input_size + 1))))
theta2 = np.matrix(np.reshape(params[hidden_size * (input_size + 1):], (num_labels, (hidden_size + 1))))
print(theta1.shape, '', theta2.shape)

```

```

(25, 401) (10, 26)

```

```

a1, z2, a2, z3, h = forward_propagate(X, theta1, theta2)
print(a1.shape, z2.shape, a2.shape, z3.shape, h.shape)

```

```

(5000, 401) (5000, 25) (5000, 26) (5000, 10) (5000, 10)

```

```

print(cost(params, input_size, hidden_size, num_labels, X, y_onehot, learning_rate))

```

```

7.242497933475441

```



```
##### end of cost function logic, below is the new part #####
```

```
# backpropagation
```

```
# taking the examples one at a time
```

```
for t in range(m):
```

```
    a1t = a1[t,:] # (1, 401)
```

```
    z2t = z2[t,:] # (1, 25)
```

```
    a2t = a2[t,:] # (1, 26)
```

```
    ht = h[t,:] # (1, 10)
```

```
    yt = y[t,:] # (1, 10)
```

```
    d3t = ht - yt # (1, 10)
```

```
    z2t = np.insert(z2t, 0, values=np.ones(1)) # (1, 26)
```

```
    # d2t (1 x 26)-matrix containing deltas of the hidden layer
```

```
    ### FILL-IN
```

```
    delta1 = delta1 + (d2t[:,1:]).T * a1t # (25, 401)
```

```
    delta2 = delta2 + d3t.T * a2t # (10, 26)
```

```
delta1 = delta1 / m
```

```
delta2 = delta2 / m
```

```
# add the gradient regularization term
```

```
delta1[:,1:] = delta1[:,1:] + (theta1[:,1:] * learning_rate) / m
```

```
delta2[:,1:] = delta2[:,1:] + (theta2[:,1:] * learning_rate) / m
```

$$\delta_j^{(3)} = a_j^{(3)} - y_j$$

$$\delta^{(2)} = (\Theta^{(2)})^T \delta^{(3)} \cdot * g'(z^{(2)})$$

```
correct = [1 if a == b else 0 for (a, b) in zip(y_pred, y)]  
accuracy = (sum(map(int, correct)) / float(len(correct)))  
print('accuracy = ', accuracy*100, '%')
```

accuracy = 98.98 %



# EXERCISE 6

```
# Run this cell to import all the libraries you need for this exercise
import numpy as np
import tensorflow as tf
from scipy.io import loadmat
```

```
# Read in data
data = loadmat('ex5data1.mat')
```

```
X_data = data['X']
y_data = data['y']
print(X_data.shape, y_data.shape)
```

```
(5000, 400) (5000, 1)
```

```
# Turn y into a one-hot vector format
from sklearn.preprocessing import OneHotEncoder
encoder = OneHotEncoder(sparse=False)
y_onehot = encoder.fit_transform(y_data)
```

```
# print one of the y[0] out
print(y_data[0], y_onehot[0,:])
```

```
[10] [0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
```

---


```
# FILL-IN: Define input x and corresponding correct answer y_
```

```
# FILL-IN: Define W1, b1, y1 of layer 1
```

```
# FILL-IN: Define W2, b2, y2 of layer 2
```

```
# FILL-IN: Define output y
```

```
# FILL-IN: Define cross_entropy and train_step
```



```
# Initializing the variables
init = tf.global_variables_initializer()

# Launch the graph
with tf.Session() as sess:
    sess.run(init)

    for _ in range(10000):
        sess.run(train_step, feed_dict={x: X_data, y_: y_onehot})

    correct_prediction = tf.equal(tf.argmax(y,1), tf.argmax(y_,1))
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
    print('accuracy: ', sess.run(accuracy, feed_dict={x: X_data, y_: y_onehot})*100, '%')
```

accuracy: 99.1599977016449 %

CONGRATULATIONS - YOU ARE DONE WITH Exercise 6 !!!